

COMP2521 19T0

Week 2, Tuesday: Algorithms!

Jashank Jeremy

jashank.jeremy@unsw.edu.au

algorithm analysis
complexity
recursion

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

Complexity

Problems, Algorithms, Programs, Processes

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

- problem something that needs to be solved
- algorithm well-defined instructions to solve the problem
- program implementation of the algorithm
in a particular programming language
- process an instance of a program being executed

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

What makes software 'good'?

correctness returns expected result for all valid inputs

robustness behaves 'sensibly' for non-valid inputs

efficiency returns results reasonably quickly (even for large inputs)

clarity clear code, easy to maintain/modify

consistency interface is clear and consistent (API or GUI)

lecture 2: correctness.

today: **efficiency**.

- algorithm runtime tends to be a function of input size
- often difficult to determine the average run time
- we tend to focus on asymptotic worst-case execution time
 - ... easier to analyse!
 - ... crucial to many applications: finance, robotics, games, ...

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

By far, the *most important* determinant of a program's efficiency.

Small, often constant-factor speedups from

- operating systems,
- compilers,
- hardware,
- implementation details

More important: an **efficient algorithm**.

Design

- complexity theory!

Implementation and Testing

- measure its properties!
 - ...run-time using *time(1)*
 - ...profiling tools like *gprof(1)*
 - ...performance counters like *pmc(3)*, *hwpmc(4)*

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

- 1 Write a program that implements an algorithm.
... which may not always be possible!
- 2 Run the program with inputs of varying size and composition.
... which may not always be possible!
... choosing good inputs is *extremely* important
- 3 Measure the actual runtime.
... which may not always be possible (or easy)!
... similar runtime environments required
- 4 Plot the results.
(Optionally, be confused about the results.)

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

- Don't necessarily use an implementation!
... Use pseudocode or something close to it.
- Characterise efficiency as a function of inputs.
- Take into account *all possible* inputs
- Generally produces a value that is environment-agnostic
... allowing us to evaluate comparative efficiency of algorithms

Absolute times will differ
between machines, between languages
...so we're not interested in absolute time.

We are interested in the *relative* change
as the problem size increases

We can use the *time(1)* command to measure execution time
(and several other interesting properties).

There are two common implementations:
one built-into the shell,
and one at `/usr/bin/time`
both are OK for our purposes.

```
$ time ./prog
```

```
./prog 0.01s user 0.02s system 97% cpu 0.028 total  
0k shared 0k local 11k max 0+3280 faults  
13+0 in 0+0+0 out 4 vcs 4 ivcs
```

Most of this information isn't interesting to us.
The *user* time is!

Redirect input into your program:

```
$ time ./prog < input > /dev/null  
$ ./mkinput | time ./prog > /dev/null
```

Time a linear search with different-sized inputs —

```
$ ./gen 100 A | time ./linear > /dev/null  
$ ./gen 1000 A | time ./linear > /dev/null
```

(repeat a number of times and average)

What is the relation between *input size* and *user time*?

If I know my algorithm is quadratic,
and I know that for a dataset of 1000 items,
it takes 1.2 seconds to run ...

- how long for 2000? **4.8 seconds**
- how long for 10,000? **120 seconds** (2 mins)
- how long for 100,000? **12000 seconds** (3.3 hours)
- how long for 1,000,000? **1200000 seconds** (13.9 days)

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

Given an array a of n elements,
where for any pair of indices i, j ,
 $i \leq j < n$ implies $a[i] \leq a[j]$
search for an element e in the array.

```
int a[N];          // array with N items
bool found = 0;
bool finished = false;
int i = 0;
while ((i < N) && (! found)) {&& (! finished)) {
    found = (a[i] == e);
    finished = (e < a[i]);
    i++;
}
```

How many comparisons do we need
for an array of size N ?

Best case: $t(N) \sim O(1)$

Worst case: $t(N) \sim O(N)$

Average case: $t(N) \sim O(N/2)$ $O(N)$

Still a *linear* algorithm!
Can we do better?

Let's start in the **middle**.

- If $e == a[N/2]$, we found e ; we're done!
- Otherwise, we split the array:
 - ... if $e < a[N/2]$, we search the left half ($a[0]$ to $a[(N/2) - 1]$)
 - ... if $e > a[N/2]$, we search the right half ($a[(N/2) + 1]$ to $a[N - 1]$)

How many comparisons do we need
for an array of size N ?

Best case: $t(n) \sim O(1)$

Worst case:

$$t(N) = 1 + t\left(\frac{N}{2}\right)$$

$$t(N) = \log_2 N + 1$$

$$t(N) \sim O(\log N)$$

In C, a line of code can do *lots* of things!

We're interested in 'primitive operations', though:
operations that can **execute in one step**,
which we can think of as hardware instructions.

(In COMP1521, we use the MIPS instruction set;
we get a feel for the primitive nature of instructions.)

Our cost-modelling will roughly follow the same lines,
but strictly we don't need to consider how long a primop takes.
We'll see why in a moment.

We express complexity using a range of *complexity models* and *complexity classes*.

Most commonly, **time complexity**,
for which we use Big-O notation,
representing asymptotic worst-case time complexity.
I'll sometimes call this WCET.

Sometimes, **space complexity** too.
(Not so much in this course, but useful!)

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

```
1  ssize_t lsearch (int a[], size_t n, int key)
2  {
3      for (size_t i = 0; i < n; i++)
4          if (a[i] == key)
5              return i;
6      return -1;
7  }
```

$$1 + (n + 1) + n(1 + 2 + \dots + 0) + 1$$

- When does the worst case occur? ... $\text{key} \notin a$
- How many data comparisons were made? ... n
- What is the worst-case cost? ... $3 + 4n$... $O(n)$

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

Growth rate is not affected (much, usually)
by constant factors or lower-order terms ...
so we discard them.

$3 + 4n$ becomes $O(n)$ — a *linear* function
 $3 + 4n + 3n^2$ becomes $O(n^2)$ — a *quadratic* function

These are an *intrinsic property* of the algorithm.

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

If a is time taken by the fastest primitive operation,
and b is time taken by the slowest primitive operation,
and $t(n)$ is the WCET of our algorithm ...

$$a \cdot (3 + 4n) \leq t(n) \leq b \cdot (3 + 4n)$$

Where does the log-base go?

$$O(\log_2 n) \equiv O(\log_3 n) \equiv \dots$$

(since $\log_b(a) \times \log_a(n) = \log_b(n)$)

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

$f(n)$ is $O(g(n))$

if $f(n)$ is asymptotically **less than or equal to** $g(n)$

$f(n)$ is $\Omega(g(n))$

if $f(n)$ is asymptotically **greater than or equal to** $g(n)$

$f(n)$ is $\Theta(g(n))$

if $f(n)$ is asymptotically **equal to** $g(n)$

Given $f(n)$ and $g(n)$, we say $f(n)$ is $O(g(n))$

if we have positive constants c and n_0 such that

$$\forall n \geq n_0, f(n) \leq c \cdot g(n)$$

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

- constant $O(1)$...constant-time execution, independent of the input size.
- logarithmic $O(\log n)$...some divide-and-conquer algorithms with trivial split/recombine operations
- linear $O(n)$...every element of the input has to be processed (in a straightforward way)
- n-log-n $O(n \log n)$...divide-and-conquer algorithms, where split/recombine is proportional to input
- quadratic $O(n^2)$...compute every input with every other input
...problematic for large inputs!
- cubic $O(n^3)$... misery
- factorial $O(n!)$... real misery
- exponential $O(2^n)$... running forever is fine, right?

Complexity

Determining

Timing

bsearch

Big-O

Theory

Recursion

tractable have a polynomial-time ('P') algorithm

... polynomial worst-case performance (e.g., $O(n^2)$)

... (useful and usable in practical applications)

intractable no tractable algorithm exists (usually 'NP'¹)

... worse than polynomial performance (e.g., $O(2^n)$)

... (feasible only for small n)

non-computable no algorithm exists (or can exist)

¹nondeterministic polynomial time, on a theoretical Turing Machine

What would be the time complexity of
inserting an element at the beginning of

... a linked list?

... an array?

What about the end?

What if it's ordered?

Recursion

Sometimes, problems can be expressed in terms of a simpler instance of the same problem.

- $1! = 1$
- $2! = 2 \times 1$
- $3! = 3 \times 2 \times 1$
- \vdots
- $(n-1)! = (n-1) \times \cdots \times 3 \times 2 \times 1$
- $(n)! = n \times (n-1) \times \cdots \times 3 \times 2 \times 1$

$$n! = (n-1)! \times n$$

Solving problems recursively in a program involves developing a program that calls itself.

base case (or *stopping case*)
no recursive call is needed

recursive case
calls the function on a smaller version of the problem

```
int factorial (int n) {  
    int result = 1;  
    for (int i = 1; i <= n; i++)  
        result *= i;  
    return result;  
}
```

```
int factorial (int n) {  
    if (n == 1) return 1;  
    else return n * factorial (n - 1);  
}
```

Recursive code can be horribly inefficient!

2^n calls is $O(k^n)$ time — exponential!

```
switch (n) {  
  case 0:  return 0;  
  case 1:  return 1;  
  default: return fib (n - 1) + fib (n - 2);  
}
```