

COMP2521 19T0

Week 5, Tuesday: Graphic Content (IV)!

Jashank Jeremy

jashank.jeremy@unsw.edu.au

weighted graphs

directed graphs

Administrivia

prac exam #1 **10 January**
at 10am, see WebCMS3 for details
(probably) no sample questions released

census date **13 January**
if you hate me and/or the course
prac exam marks back before then

assignment 2 part 1 is out now:
the Fury of Dracula: the View
make sure you have a group on WebCMS 3

Assignment 2, Part 1

the Fury of Dracula: the View

use a version control system
like Fossil, Git, SVN, etc.

use documentation tools like Doxygen

start sooner rather than later;
write some tests before you begin

Digraphs

Applications
Terminology
Representation
DAGs

Wgraphs

Directed Graphs

Directed Graphs

Digraphs

Applications
Terminology
Representation
DAGs

Wgraphs

We've mostly considered *undirected* graphs:
an edge relates two vertices equivalently.

Some applications require us to consider
directional edges: $v \rightarrow w \neq w \rightarrow v$
e.g., 'follow' on Twitter, one-way streets, etc.

In an **directed graph** or **digraph**:
edges have direction;
self-loops are allowed;
'parallel' edges are allowed.


Directed Graphs

Example

Digraphs

Applications
Terminology
Representation
DAGs

Wgraphs



Where can we get to from g ?
Can we get to e from anywhere else?

domain	vertex is...	edge is...
WWW	web page	hyperlink
chess	board state	legal move
scheduling	task	precedence
program	function	function call
journals	article	citation

- Is there a directed path from s to t ? (**transitive closure**)
- What is the shortest path from s to t ? (**shortest path search**)
- Are all vertices mutually reachable? (**strong connectivity**)
- How can I organise a set of tasks? (**topological sort**)
- How can I crawl the web? (**graph traversal**)
- Which web pages are important? (**PageRank**)

in-degree or $d^{-1}(v)$: the number of directed edges leading **into** a vertex
out-degree or $d(v)$: the number of directed edges leading **out of** a vertex

sink a vertex with out-degree 0;

source a vertex with in-degree 0

reachability indicates existence of directed path:
if a directed path v, \dots, w exists,
 w is reachable from v

strongly connected indicates mutual reachability:
if both paths v, \dots, w and w, \dots, v exist,
 v and w are strongly connected


strong connectivity every vertex reachable from every other vertex;

strongly-connected component maximal strongly-connected subgraph

Similar choices as for undirected graphs:


- adjacency matrix ... asymmetric, sparse; less space efficient
- adjacency lists ... fairly common solution
- edge lists ... order of edge components matters
- linked data structures ... pointers inherently directional

Can we make our undirected graph implementations directed? Yes!



$$\begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 \end{bmatrix}$$

unweighted, undirected




$$\begin{bmatrix} - & 1 & 1 & 1 & - \\ - & - & 1 & - & - \\ - & - & - & 1 & 1 \\ - & - & - & - & - \\ - & - & 1 & 1 & - \end{bmatrix}$$

unweighted, **directed**

	storage	edge add	has edge	outdegree
adj.matrix	$O(V + V^2)$	$O(1)$	$O(1)$	$O(V)$
adj.list	$O(V + E)$	$O(d(v))$	$O(d(v))$	$O(d(v))$

Overall, adjacency lists tend to be ideal:
real digraphs tend to be sparse
(large V , small average $d(v)$);
algorithms often iterate over v 's edges

Directed Acyclic Graphs



Is it a tree? Is it a graph?
No: it's a DAG, a directed acyclic graph.

Tree-like: each vertex has 'children'.
Graph-like: a child vertex may have multiple parents.

Directed Acyclic Graphs

Application: the Topological Sort

NOT EXAMINABLE (and not taught until '4128)

The most common application of a DAG is *topological sorting*:
ordering vertices such that, for any vertices u and v ,
if u has a directed edge to v , then v comes after u in the ordering.

Computable with a DFS, tracking *post-order sequence*:
vertices only added after their children have been visited
 \Rightarrow a valid topological ordering

dependency problems: *make(1)*, spreadsheets
version-control systems: Git, Fossil, etc.

Mostly the same algorithms as for undirected graphs:
DFS and BFS should all Just Work

e.g., Web crawling: visit every page on the web.
BFS with implicit graph;
on visit, scans page for content, keywords, links
... assumption: www is fully connected.

Weighted Graphs

Weighted Graphs

Some applications require us to consider a **cost** or **weight** assigned to a relation between two nodes.

Often, we use a geometric interpretation:
low weight \Rightarrow short edge;
high weight \Rightarrow long edge;

Weights aren't always geometric:
some weights are **negative**.
(We assume we have non-negative weights,
as graphs with negative weights tend to cause problems...)

Adjacency matrix:

- store *weight* in each cell, not just true/false.
- need some “no edge exists” value: zero might be a valid weight.

Adjacency list

- add weight to each list node

Edge list:

- add weight to each edge


Linked data structure:

- links become link/weight pairs

Works for directed and undirected graphs!


Weighted Graphs

Implementation: Adjacency Matrix



$$\begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 1 \\ 1 & 0 & 1 & 0 & 1 \\ 0 & 0 & 1 & 1 & 0 \end{bmatrix}$$

unweighted, undirected



$$\begin{bmatrix} - & 0.2 & 0.4 & 0.5 & - \\ 0.2 & - & 0.5 & - & - \\ 0.4 & 0.5 & - & 0.1 & 0.1 \\ 0.5 & - & 0.1 & - & 0.9 \\ - & - & 0.1 & 0.9 & - \end{bmatrix}$$

weighted, undirected

Weighted Graph Problems

The shortest path problem:

- find the minimum cost path between two vertices
- edges may be directed or undirected
- assuming non-negative weights!


minimum spanning trees (MST):

- find the *weight-minimal* set of edges that connect all vertices in a weighted graph
- multiple solutions may exist!
- assuming undirected, non-negatively-weighted graphs

Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs
Kruskal
Prim
Others



What's the shortest path from 0 to 3?

What's the least-hops (shortest unweighted path) from 0 to 2?

What is the minimum spanning tree?

Shortest-Path Search

Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs
Kruskal
Prim
Others

Shortest-path is useful in navigation and route-finding on physical maps, in computer networks, etc.

Several flavours of shortest-path searches exist:

- source-target** the shortest path from v to w ;
- single-source** the shortest path from v to all other vertices;
- all-pairs** the shortest paths for all pairs of v, w

Shortest-Path Search

Formally

Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs
Kruskal
Prim
Others

On graph G , the weight of p (as $\text{weight}(p)$) is the sum of weights of p 's edges.

The shortest path between v and w is a simple path $p = [v, \dots, w]$, where no other simple path $q = [v, \dots, w]$, with $q \neq p$, has a lesser weight (i.e., $\forall q, \text{weight}(p) < \text{weight}(q)$).

Assuming a weighted graph, with no negative weights. (On an unweighted graph, devolves to least-hops.)

Given a weighted graph G , and a start vertex v ,
we want shortest paths from v to all other vertices.

ASIDE how do we represent it?

we get a vertex-indexed array of distances from v ,
and a vertex-indexed array of shortest-path predecessors
... it's a spanning tree rooted at v .
(Spanning trees can have weighted and/or directed edges, too!)

```

sssp (Graph  $g$ , vertex  $v$ ):
    dists[] := [ $\infty$ , ...]
    dists[v] = 0
    pq := NEWPQUEUE
    for each  $e := (s, t, \omega)$  in ADJACENT( $v$ ),
        ENPQUEUE(pq, ( $s, t$ ),  $\omega$ )

    while LENGTH(pq) > 0:
        ( $s, t$ ),  $\omega$  := DEPQUEUE(pq)
        get edges that connect  $s$  and  $t$ 
        relax along edge if new distance is better
        add edges with total path weights
    
```

“Edge relaxation” along edge e from s to t :

$\text{dist}[s]$ is length of some path from v to s ;
 $\text{dist}[t]$ is length of some path from v to t
 if e gives shorter path v to t via s ,
 update $\text{dist}[t]$ and $\text{st}[t]$.

Relaxation updates data on t , if we find a shorter path from v .

```

if (dist[s] + e.weight < dist[t]) {
    dist[t] = dist[s] + e.weight;
    pqueue_en (pq, t, dist[t]);
    st[t] = s;
}
    
```

COMP2521
19T0 lec08

cs2521@
jashankj@

Single-Source Shortest-Path Search

Demonstration

Digraphs

Wgraphs

Shortest Paths

Single-Source, Dijkstra

Single-Source, Others

All-Pairs

MSTs

Kruskal

Prim

Others

	0	1	2	3	4	5	6	7	8
dist									
st									

COMP2521
19T0 lec08

cs2521@
jashankj@

Single-Source Shortest-Path Search

Results; Complexity

Digraphs

Wgraphs

Shortest Paths

Single-Source, Dijkstra

Single-Source, Others

All-Pairs

MSTs

Kruskal

Prim

Others

Once this algorithm has run:
shortest path distances are in `dist`;
predecessors in `st` array; trace for a path

COMPLEXITY:
using an adjacency list and a heap: $O(E \log V)$;
using an adjacency matrix: $O(V^2)$.

Just a graph traversal (*a la* BFS, DFS),
but using a PQueue, instead of a Stack/Queue.

This algorithm is usually known as
Dijkstra's algorithm.
Sedgewick calls this a PRIORITY-FIRST SEARCH.

COMP2521
19T0 lec08

cs2521@
jashankj@

Single-Source Shortest-Path Search

Situation Overview

Digraphs

Wgraphs

Shortest Paths

Single-Source, Dijkstra

Single-Source, Others

All-Pairs

MSTs

Kruskal

Prim

Others

constraint	algorithm	cost	remark
single-source shortest:			
non-negative weights	Dijkstra	V^2	optimal (dense)
non-negative weights	Dijkstra	$E \log V$	conservatively
acyclic	source-queue	E	optimal
no negative cycles	Bellman-Ford	VE	improvements?
(none)	?	?	NP-hard

Do Dijkstra's sssp at every vertex.
(This sucks as much as it sounds like it does.)

Floyd-Warshall.
(Out of scope, see '4121/'4128).

constraint	algorithm	cost	remark
all-pairs shortest:			
non-negative weights	Floyd	V^3	same for all
non-negative weights	Dijkstra (PFS)	$VE \log V$	conservatively
acyclic	DFS	VE	same for all
no negative cycles	Floyd	V^3	same for all
no negative cycles	Johnson	$VE \log V$	conservatively
(none)	?	?	NP-hard

originally, Otakar Borůvka in 1926:
most economical construction of electric power network
(*O jistém problému minimálním*, 'On a certain minimal problem')

routing and network layout:
electricity, telecommunications, electronic, road, ...
widely applicable \Rightarrow intensely studied problem

Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs

MSTs

Kruskal
Prim
Others

A spanning tree ST of a graph $G(V, E)$ is a subgraph $G'(V, E')$, such that $E' \subseteq E$.
ST is connected (spanning) and acyclic (tree)

A minimum spanning tree MST of a graph G is a spanning tree of G , where the sum of edge weight is no larger than any other spanning tree.

So: how do we (efficiently) find a MST for G ?

Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs

Kruskal
Prim
Others

Kruskal's Algorithm

take all edges, sorted according to their weight;
then, for each edge:
add it to the proto-MST;
unless it would introduce a cycle,

Cycle-checking is really expensive
(DFS everything!)
Sedgewick has a 'union-find' that works fine here

Sorting dominates overall:
 $O(E \log E)$.

Digraphs


Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs

Kruskal
Prim
Others

Kruskal's Algorithm

Demonstration (I)



v	0	0	0	1	1	2	2	3	3	4	4	4	7
w	1	3	4	2	4	4	5	4	6	5	6	7	8
ω	3	4	12	5	10	13	1	6	2	7	11	8	9


Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs

Kruskal

Prim
Others



v	2	3	0	0	1	3	4	4	7	1	4	0	2
w	5	6	1	3	2	4	5	7	8	4	6	4	4
ω	1	2	3	4	5	6	7	8	9	10	11	12	13

Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs

Kruskal

Prim
Others

Prim-Jarník-Dijkstra Algorithm

Another approach to computing an MST for graph $G(V, E)$ discovered by Prim (1957), Jarník (1930), Dijkstra (1959)

- 1 start from any vertex s and with an empty MST
- 2 choose edge not already in MST to add
 - must not contain a self-loop
 - must connect to a vertex already on MST (on the *fringe*)
 - must have minimal weight of all such edges
- 3 check to see whether adding the new edge brought any of the non-tree vertices closer to the MST
- 4 repeat until MST covers all vertices

basically just Dijkstra's sssp algorithm,
just a graph search but using a PQueue;
 $O(E \log V)$ (adjacency lists, heap)
or $O(V^2)$ (adjacency matrix).


Digraphs

Wgraphs

Shortest Paths
Single-Source, Dijkstra
Single-Source, Others
All-Pairs
MSTs

Kruskal

Prim
Others



v	0	0	0	1	1	2	2	3	3	4	4	4	7
w	1	3	4	2	4	4	5	4	6	5	6	7	8
ω	3	4	12	5	10	13	1	6	2	7	11	8	9

Digraphs

Wgraphs

Shortest Paths

Single-Source,
Dijkstra

Single-Source,
Others

All-Pairs

MSTs

Kruskal

Prim

Others

- Kruskal: grow many forests
- Prim/Jarník/Dijkstra: maintain connectivity on frontier
- Borůvka/Sollin: component-wise
- Tarjan/Karger/Klein: randomised
- Chazelle: deterministic; best-performing