

COMP2521 19T0

Week 7, Tuesday: A Question of Balance

Jashank Jeremy

jashank.jeremy@unsw.edu.au

radix sort
balanced trees

Sorting

Radix Sorting

Can we decompose our keys?
Radix sorts let us deal with this case.

Keys are values in some base- R number system.

e.g., binary, $R = 2$; decimal, $R = 10$;

ASCII, $R = 128$ or $R = 256$; Unicode, $R = 2^{16}$

Sorting individually on each part of the key at a time:
digit-by-digit, character-by-character, rune-by-rune, *etc.*

Radix Sorting, Most-Significant-Digit First

Consider characters, digits, bits, runes, etc.,
from **left to right**;
partitioning input into R pieces according to key $\cdot 0$;
recurse into each piece, using successive keys —
key $\cdot 1$, key $\cdot 2$, ..., key $\cdot w$

1019	2301	3129	2122
1019	2301	3129	2122
1019	2301	2122	3129
1019	2122	2301	3129

with $R = 2$, roughly a quicksort.

Radix Sorting, Least-Significant-Digit First

Consider characters, digits, bits, runes, etc.,
from **right to left**;
use a **stable** sort using the d th digit as key,
using (e.g.,) key-indexed counting sort.

1019	2301	3129	2122
1019	2301	3129	2122
2301	2122	1019	3129
2301	1019	2122	3129
1019	2122	3129	2301
1019	2122	2301	3129

this *will not work* if the sort is not stable!

Radix Sort

Analysis; Summary

Complexity: $O(w(n + R)) \approx O(n)$,
where w is the 'width' of data;
the algorithm makes w passes over n keys

LSD

Not in-place: $O(n + R)$ extra space required.
May be stable! Usable on variable length data.

MSD

Not in-place: $O(n + DR)$ extra space required.
(D is the recursion depth.)
May be stable! Usable on variable length data.
Can complete *before* examining all of all keys.

Balanced Trees

Recap: The Search Problem

input a key value
output item(s) containing that key

Common variations:

- keys are unique; key value matches 0 or 1 items
- multiple keys in search, items containing any key
- multiple keys in search/item, items containing all keys

We assume: keys are unique, each item has one key.

Recap: Trees; Binary Trees

Trees are branched data structures,
consisting of **nodes** and **edges**, with no cycles.

Each node contains a value.
Each node has edges to $\leq k$ other nodes.
For now, $k = 2$ — binary trees

Trees can be viewed as a set of nested structures:
each node has k (possibly empty) **subtrees**.

Recap: Binary Search Trees

For all nodes in the tree:
the values in the **left** subtree
are **less than** the node value the

values in the **right** subtree
are **greater than** the node value

A binary tree of n nodes
is **degenerate** if its height is
at most $n - 1$.


A binary tree of n nodes
is **balanced** if its height is
at least $\lfloor \log_2 n \rfloor$.

Structure tends to be determined
by order of insertion:
[4, 2, 1, 3, 6, 5, 7] vs [6, 5, 2, 1, 3, 4, 7]

Binary Search Trees

The Worst Case

Ascending-ordered or
descending-ordered data
is a **pathological case**:
we always right- or left-insert
along the spine of the tree.



Binary Search Trees

Performance

Cost for **insertion**:
balanced $O(\log_2 n)$, degenerate $O(n)$
(we always traverse the height of the tree)


Cost for **search/deletion**:
balanced $O(\log_2 n)$, degenerate $O(n)$
(worst case, key $\notin \tau$; traverse the height)

We want to build balanced trees.

PERFECTLY BALANCED
a *weight-balanced* or
size-balanced tree has,
for every node,
 $|\text{SIZE}(l) - \text{SIZE}(r)| < 2$

LESS STRINGENTLY
a *height-balanced* tree has,
for every node,
 $|\text{HEIGHT}(l) - \text{HEIGHT}(r)| < 2$

Balanced or Not?
(I)




$$\begin{aligned}\text{SIZE}(\tau_4) &= 5 \\ \text{SIZE}(\tau_2) &= 3 \\ \text{SIZE}(\tau_5) &= 1 \\ \text{SIZE}(\tau_1) &= 1 \\ \text{SIZE}(\tau_3) &= 1 \\ \text{SIZE}(\tau_2) - \text{SIZE}(\tau_5) &= 2\end{aligned}$$

NOT SIZE BALANCED

$$\begin{aligned}\text{HEIGHT}(\tau_4) &= 2 \\ \text{HEIGHT}(\tau_2) &= 1 \\ \text{HEIGHT}(\tau_5) &= 0 \\ \text{HEIGHT}(\tau_1) &= 0 \\ \text{HEIGHT}(\tau_3) &= 0 \\ \text{HEIGHT}(\tau_2) - \text{HEIGHT}(\tau_5) &= 1\end{aligned}$$

HEIGHT BALANCED

Balanced or Not?
(II)




$$\begin{aligned}\text{SIZE}(\tau_4) &= 5 \\ \text{SIZE}(\tau_2) &= 2 \\ \text{SIZE}(\tau_5) &= 2 \\ \text{SIZE}(\tau_3) &= 1 \\ \text{SIZE}(\tau_6) &= 1\end{aligned}$$

SIZE BALANCED

$$\begin{aligned}\text{HEIGHT}(\tau_4) &= 2 \\ \text{HEIGHT}(\tau_2) &= 1 \\ \text{HEIGHT}(\tau_5) &= 1 \\ \text{HEIGHT}(\tau_3) &= 0 \\ \text{HEIGHT}(\tau_6) &= 0\end{aligned}$$

HEIGHT BALANCED



Let's look at τ_3 .


$SIZE(\tau_2) = 2$
 $SIZE(\tau_\emptyset) = 0$
 $2 - 0 = 2 \not< 2$

NOT SIZE BALANCED

Let's look at τ_5 .

$HEIGHT(\tau_\emptyset) = 0$
 $HEIGHT(\tau_6) = 1$
 $|0 - 1| = 1 < 2$

HEIGHT BALANCED



Let's look at τ_4 .

$SIZE(\tau_2) = 3$
 $SIZE(\tau_\emptyset) = 0$
 $3 - 0 = 3 \not< 2$

NOT SIZE BALANCED


Let's look at τ_4 .

$HEIGHT(\tau_2) = 1$ $HEIGHT(\tau_\emptyset) = 0$
 $1 - 0 = 1 < 2$

HEIGHT BALANCED

LEFT ROTATION and RIGHT ROTATION:

a pair of ‘primitive’ operations
that change the balance of a tree
whilst maintaining a search tree.



Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives

Rotation

Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

```
btree_node *btree_rotate_right (btree_node *n1)
{
    if (n1 == NULL) return NULL;
    btree_node *n2 = n1->left;
    if (n2 == NULL) return n1;
    n1->left = n2->right;
    n2->right = n1;
    return n2;
}
```

n_1 starts as the root of this subtree and is demoted;
 n_2 starts as the left subtree of this tree, and is promoted.

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives

Rotation

Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

```
btree_node *btree_rotate_left (btree_node *n2)
{
    if (n2 == NULL) return NULL;
    btree_node *n1 = n2->right;
    if (n1 == NULL) return n2;
    n2->right = n1->left;
    n1->left = n2;
    return n1;
}
```

n_2 starts as the root of this subtree and is demoted;
 n_1 starts as the right subtree of this tree, and is promoted.

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives

Rotation

Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

A way to brute-force some balance into a tree:
lifting some k th index to the root.

PARTITION :: BTree \rightarrow Word \rightarrow BTree

PARTITION Empty k = Empty

PARTITION (Node n l r) k

| $k < \text{SIZE } l$ = ROTATER (Node n (PARTITION l k) r)

| $\text{SIZE } l < k$ = ROTATEL (Node n l (PARTITION r ($k - 1 - \text{SIZE } l$)))

| otherwise = Node n l r


Sorting

Balanced
Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation

Partition

Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay



What happens if we partition at index 3 (node 4)?

Sorting

Balanced
Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation

Partition

Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

```
btree_node *btree_partition (btree_node *tree, size_t k)
{
    if (tree == NULL) return NULL;
    size_t lsize = size (tree->left);
    if (lsize > k) {
        tree->left = btree_partition (tree->left, k);
        tree = btree_rotate_right (tree);
    }
    if (lsize < k) {
        tree->right = btree_partition (tree->right, k - 1 - lsize);
        tree = btree_rotate_left (tree);
    }
    return tree;
}
```

Sorting

Balanced
Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation

Partition

Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

With our primitive operations in hand —

$\text{ROTATE_L} :: \text{BTree} \rightarrow \text{BTree}$
 $\text{ROTATE_R} :: \text{BTree} \rightarrow \text{BTree}$
 $\text{PARTITION} :: \text{BTree} \rightarrow \text{Word} \rightarrow \text{BTree}$

— let's go balance some trees!

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees

Properties
Primitives
Rotation
Partition

Simple Approaches

Global

Root Insert
Random Trees
Complex Approaches
Splay

Move the median node to the root,
by partitioning on $\text{SIZE } \tau/2$;
then, balance the left subtree,
and balance the right subtree.

```
btree_node *btree_balance_global (btree_node *tree)
{
    if (tree == NULL) return NULL;
    if (size (tree) < 2) return tree;
    tree = partition (tree, size (tree) / 2);
    tree->left = btree_balance_global (tree->left);
    tree->right = btree_balance_global (tree->right);
    return tree;
}
```

Approach #1: Global Rebalancing

Problems

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees

Properties
Primitives
Rotation
Partition

Simple Approaches

Global

Root Insert
Random Trees
Complex Approaches
Splay

- cost of rebalancing:
for many trees, $O(n)$; for degenerate trees, $O(n \log n)$
- what if we insert more keys?
 - rebalance on every insertion
 - rebalance every k insertions; what k is good?
 - rebalance when imbalance exceeds threshold.

we either have more costly insertions
or degraded performance for (possibly unbounded) periods.
... given a sufficiently dynamic tree, sadness.

Global vs Local Rebalancing

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees

Properties
Primitives
Rotation
Partition

Simple Approaches

Global

Root Insert
Random Trees
Complex Approaches
Splay

GLOBAL REBALANCING

walks every node, balances its subtree;
 \Rightarrow perfectly balanced tree — at cost.

LOCAL REBALANCING

do small, incremental operations
to improve the overall balance of the tree
... at the cost of imperfect balance

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

amortisation: do (a small amount) more work now to avoid more work later
randomisation: use randomness to reduce impact of BST worst cases
optimisation: maintain structural information for performance

Root Insertion

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

How do we insert a node at the root of a tree?
(Without having to rearrange all the nodes?)

We do a leaf insertion ...
... and rotate the new node up the tree.

More work? **No!**
Same complexity as leaf insertion,
but more actual work is done: **amortisation**.

(Side-effect: recently-inserted items are close to the root.
Depending on what you're doing, this might be very useful!)

Root Insertion

C Implementation

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

```
btree_node *btree_insert_root (btree_node *tree, Item it)
{
    if (tree == NULL)
        return btree_node_new (it, NULL, NULL);
    if (less (it, tree->value)) {
        tree->left = btree_insert_root (tree->left, it);
        tree = btree_rotate_right (tree);
    } else {
        tree->right = btree_insert_root (tree->right, it);
        tree = btree_rotate_left (tree);
    }
    return tree;
}
```

BSTs don't have control over insertion order.
worst cases — (partially) ordered data — are common.

to minimise the likelihood of a degenerate tree,
we randomly choose which level to insert a node;
at each level, probability depends on remaining tree size.

do a 'normal' leaf insertion, most of the time.
randomly (with a certain probability),
do a root insertion of a value.

```
btree_node *btree_insert_rand (btree_node *tree, Item it)
{
    if (tree == NULL)
        return btree_node_new (it, NULL, NULL);
    if (rand () < (RAND_MAX / size (tree)))
        return btree_insert_root (tree, it);
    else if (less (it, tree->value))
        tree->left = btree_insert_rand (tree->left, it);
    else
        tree->right = btree_insert_rand (tree->right, it);
    return tree;
}
```

building a randomised BST is equivalent to
building a standard BST with
a random initial permutation of keys.

worst-case, best-case, average-case performance:
same as a standard BST —
but with no penalty for ordering!

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert

Random Trees

Complex Approaches
Splay

We could do something similar for deletion:
when choosing a node to promote,
choose randomly from the
in-order predecessor or successor

Splay Trees

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

Root insertion can still leave us
with a degenerate tree.

Splay trees vary root-insertion,
by considering *three* levels of the tree
— parent, child, grandchild —
and performing double-rotations based on p-c-g orientation;
the idea: double-rotations improve balance.

No guarantees, but *improved* performance.

“... their performance is amortised by
the amount of effort required to understand them.”
— me, 2016

Splay Trees


Double-Rotation Cases

Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

Four choices to consider for a double-rotation:




Sorting

Balanced
Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches

Splay

ROTATER τ_M ROTATER τ_E



Splay Rotations

Double-Rotation: Left, Right


Sorting

Balanced
Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches

Splay

ROTATEL τ_E ROTATER τ_M



Splay Rotations

Double-Rotation: Right, Left


Sorting

Balanced
Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches

Splay

ROTATER τ_Q ROTATEL τ_M



Sorting

Balanced Trees

Recap
Searching
Trees, B-Trees
Search Trees
Properties
Primitives
Rotation
Partition
Simple Approaches
Global
Root Insert
Random Trees
Complex Approaches
Splay

ROTATE \bar{L} τ_M

ROTATE \bar{L} τ_Q

