

# COMP2521 19T0

## Week 6, Thursday: Order! Order (II)

Jashank Jeremy

jashank.jeremy@unsw.edu.au

more sorting algorithms  
non-comparing sorts

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

# Sorting

## Sorting

### Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

divide-and-conquer algorithms  
break up, or **shard**, the problem  
into (easier) computations on smaller pieces,  
and **combine** the results.

(usually) easy to implement recursively!  
(usually) easy to implement in parallel!

## Sorting

### Divide-and-Conquer

#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap

- 1 If a collection has less than two elements, it's sorted. Otherwise, split it into 2 parts.
- 2 Sort both parts separately.
- 3 Combine the sorted collections to return the final result.

Copy elements from the inputs one at a time,  
giving preference to the smaller of the two.  
When one list is empty,  
copy the rest of the elements from the other.

```
/// In-place merge on array `a`,  
/// of slices `a[lo..mid-1]` and `a[mid..hi]`.  
void merge (Item a[], size_t lo, size_t mid, size_t hi);
```

```
/// Out-of-place merge, into `dest[0..ndest-1]`,  
/// from `a[0..nA-1]` and `b[0..nB-1]`.  
void merge (Item dest[], size_t ndest,  
            Item a[], size_t nA, Item b[], size_t nB);
```

A divide-and-conquer sort:

**partition** the input into two equal-sized parts.

**recursively** sort each of the partitions.

**merge** the two now-sorted partitions back together.

## Sorting

### Divide-and-Conquer


#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap



## Sorting

### Divide-and-Conquer


#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap





### Sorting

#### Divide-and-Conquer

#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap

```
void sort_merge (Item a[], size_t lo, size_t hi)
{
    if (hi <= lo) return;
    size_t mid = (lo + hi) / 2;
    sort_merge (a, lo, mid);
    sort_merge (a, mid+1, hi);
    merge (a, lo, mid, hi);
}
```

```
void merge (Item a[], size_t lo, size_t mid, size_t hi)
{
    Item *tmp = calloc (hi - lo + 1, sizeof (Item));
    size_t i = lo, j = mid + 1, k = 0;

    // Scan both segments, copying to `tmp'.
    while (i <= mid && j <= hi)
        tmp[k++] = less (a[i], a[j]) ? a[i++] : a[j++];

    // Copy items from unfinished segment.
    while (i <= mid) tmp[k++] = a[i++];
    while (j <= hi) tmp[k++] = a[j++];

    // Copy `tmp' back to main array.
    for (i = lo, k = 0; i <= hi; a[i++] = tmp[k++]);

    free (tmp);
}
```

How many steps does it take  
to sort a collection of  $N$  elements?

Splitting arrays into two halves: constant time.  
To re-combine,  $N$  steps.

$$T(N) = N + 2T(N/2)$$

substitute  $N := 2^N$ ; then:

$$T(2^N) = 2^N + 2T(2^N/2)$$

$$T(2^N) = 2^N + 2T(2^{N-1})$$

## Sorting

## Divide-and-Conquer

## Merge

## Quick

## Non-Comparison

## Key-Indexed

## Heap

divide out  $2^N$ ; then:

$$\begin{aligned}T(2^N) / 2^N &= 1 + 2 T(2^{N-1}) / (2^N) \\T(2^N) / 2^N &= 1 + T(2^{N-1}) / (2^{N-1})\end{aligned}$$

expanding, we get:

$$\begin{aligned}&1 + (1 + T(2^{N-2}) / (2^{N-2})) \\&1 + (1 + (1 + T(2^{N-3}) / (2^{N-3}))) \\&\dots = N\end{aligned}$$

$$\begin{aligned}T(2^N) / 2^N &= N \\T(2^N) &= 2^N N\end{aligned}$$

$$T(N) = N \log_2 N$$

## Sorting

### Divide-and-Conquer

#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap

How many steps does it take  
to sort a collection of  $N$  elements?

- split array into equal-sized partitions  
halving at each level  $\Rightarrow \log_2 N$  levels
- same operations happen at every recursive level
- each 'level' requires  $\leq N$  comparisons —  
worst case: two arrays exactly interleaved,  $N$  comparisons

## Sorting

### Divide-and-Conquer

#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap

Merge sort is  $O(n \log n)$ .

Generally, stable...  
... as long as the merge is stable.

*Not* in-place:  
 $O(n)$  memory for merge;  $O(\log n)$  stack space.

Oblivious:  
 $O(n \log n)$  best case, average case, worst case

## Sorting

### Divide-and-Conquer

#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap

Straightforward!

- Traverses input in sequential order.
- Don't need extra space for merging list.
- Works top-down and ... bottom-up?

## Sorting

### Divide-and-Conquer

#### Merge

#### Quick

#### Non-Comparison

#### Key-Indexed

#### Heap

An approach that works non-recursively!

- on each pass, our array contains sorted *runs* of length  $m$ .
- initially,  $N$  sorted runs of length 1.
- The first pass merges adjacent elements into runs of length 2.
- The second pass merges adjacent elements into runs of length 4.
- ... continue until we have a single sorted run of length  $N$ .

Can be used for *external* sorting;  
*e.g.*, sorting disk-file contents



## Sorting

## Divide-and-Conquer

## Merge

## Quick

## Non-Comparison

## Key-Indexed

## Heap

```
#define MIN(a,b) ((a) < (b) ? (a) : (b))

void sort_merge_bu (Item a[], size_t lo, size_t hi)
{
    for (size_t m = 1; m <= lo - hi; m *= 2)
        for (size_t i = lo; i <= hi - m; i += 2 * m) {
            size_t end = MIN (i + 2*m - 1, hi);
            merge (a, i, i + m - 1, end);
        }
}
```

Merge sort uses a trivial split operation;  
all the heavy lifting is in the *merge* operation.

Can we split the collection in a more intelligent way,  
so combining the results is easier?

...e.g., making sure all elements in one part  
are less than elements in the second part?

## Sorting

Divide-and-Conquer

Merge


Quick

Non-Comparison

Key-Indexed

Heap

to partition array  $a$   
at some index  $i$  (the 'pivot'),  
we need to swap elements such that,  
for other indices  $j$  and  $k$ ,  
 $j < i$  implies  $a[j] \leq a[i]$   
 $k > i$  implies  $a[i] \leq a[k]$



Assuming we have a partition function,  
this looks very similar to merge sort.

```
void sort_quick_naive (Item a[], size_t lo, size_t hi)
{
    if (hi <= lo) return;
    size_t part = partition (a, lo, hi);
    sort_quick_naive (a, lo, part - 1);
    sort_quick_naive (a, part + 1, hi);
    // look, ma! no merge!
}
```

```
size_t partition (Item a[], size_t lo, size_t hi)
{
    Item v = a[lo]; // our 'pivot' value.
    size_t i = lo + 1, j = hi;
    for (;;) {
        while (less (a[i], v) && i < j) i++;
        while (less (v, a[j]) && i < j) j--;
        if (i == j) break;
        swap_idx (a, i, j);
    }
    j = less (a[i], v) ? i : i - 1;
    swap_idx (a, lo, j);
    return j;
}
```

Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

How many steps does it take  
to sort a collection of  $N$  elements?

$N$  steps to partition an array...  
constant-time combination of sub-results.

best-case (equal sized partitions):  $O(N \log N)$

worst-case (one part contains all elements):

$$\begin{aligned} T(N) &= N + T(N-1) = N + (N-1) + T(N-2) \\ &\dots = N(N+1)/2, \text{ which is } O(N^2) \end{aligned}$$

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

Quick sort with naïve partition is...

Unstable (in this implementation)...  
... but can be made stable.

In-place: partitioning is done in-place;  
stack depth is  $O(N)$  worst-case,  $O(\log N)$  average

Oblivious.

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

Picking the first or last element as pivot  
is **an absolutely terrible life choice**.

... existing order is a worst case.

... existing *reverse* order is a worst case.

partition always gives us parts of size  $N - 1$  and 0.

Our ideal pivot is the *median* value.

Our worst pivot is the largest/smallest value.

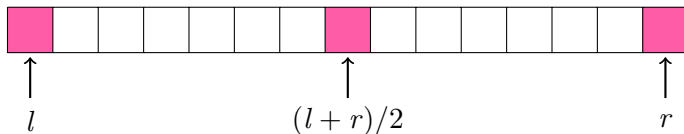
We can reduce the probability of picking a bad pivot...



# Quick Sort with Median-of-Three Partition

Pick three values: left-most, middle, right-most.  
Pick the median of these three values as our pivot.

Ordered data is no longer a worst-case scenario.  
In general, doesn't eliminate the worst-case ...  
... but makes it much less likely.



# Quick Sort with Median-of-Three Partitioning

## Sorting

Divide-and-Conquer


Merge

Quick

Non-Comparison

Key-Indexed

Heap



- 1 Pick  $a[l]$ ,  $a[r]$ ,  $a[(l + r)/2]$
- 2 Swap  $a[r - 1]$  and  $a[(l + r)/2]$
- 3 Sort  $a[l]$ ,  $a[r - 1]$ ,  $a[r]$ , such that  $a[l] \leq a[r - 1] \leq a[r]$
- 4 Partition on  $a[l + 1]$  to  $a[r - 1]$ .

# Quick Sort with Median-of-Three Partitioning

C Implementation

Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

```
void qs_median3 (Item a[], size_t lo, size_t hi)
{
    size_t mid = (lo + hi) / 2;
    if (less (a[mid], a[lo])) swap_idx (a, lo, mid);
    if (less (a[hi], a[mid])) swap_idx (a, mid, hi);
    if (less (a[mid], a[lo])) swap_idx (a, lo, mid);
    // now, we have a[lo] <= a[mid] <= a[hi]
    // swap a[mid] to a[lo+1] to use as pivot
    swap_idx (a, lo+1, mid);
}
```

```
void sort_quick_m3 (Item a[], size_t lo, size_t hi)
{
    if (hi <= lo) return;
    qs_median3 (a, lo, hi);
    size_t part = partition (a, lo + 1, hi - 1);
    sort_quick_m3 (a, lo, part - 1);
    sort_quick_m3 (a, part + 1, hi);
}
```

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

For small sequences (when  $n < 5$ , say),  
quick sort is **expensive**  
because of the recursion overhead.

With a **sub-file cutoff**, we have two choices:  
use a different algorithm on small partitions; or  
do a second sort after the quicksort finishes.  
(Insertion sort is a good choice: lots of almost-sorted data!)

### Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

For sequences with many duplicate keys,  
partitioning can screw up badly.

instead, do a **three-way** partition:

$\text{keys} < a[i], = a[i], > a[i].$

Straightforward to do...  
if we just use the first or last element as pivot  
(which means we're vulnerable to ordered data again)

using a random or median-of-three pivot  
is now  $O(n)$  not  $O(1)$

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

Design of modern CPUs mean,  
for sorting arrays in RAM  
quicksort *generally* outperforms mergesort.

quicksort is more 'cache friendly':  
good locality of access on arrays

on the other hand, mergesort is  
readily stable, readily parallel,  
more efficient with slower data;  
a good choice for sorting linked lists

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

If we have 3 items,  
then  $3! = 6$  possible permutations as input.  
( $n$  items implies  $n!$  possible permutations.)

If we do 1 comparison,  
we can form two categories (true, false).  
( $k$  comparisons implies  $2^k$  categories.)



## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

$n$  items implies  $n!$  possible permutations.  
 $k$  comparisons implies  $2^k$  categories.

We need to do enough comparisons so

$$n! \leq 2^k.$$

$$\log_2 n! \leq \log_2 2^k$$

$$\log_2 n! \leq k$$

... applying Stirling's approximation, and waving our hands:

$$n \log n < k.$$

the theoretical lower bound on  
worst-case execution time  
for comparison-based sorts is  $O(n \log n)$ .  
(Quicksort, mergesort are pretty much  
as good as it gets, for unknown data.)

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

All the sorts so far have been  
**comparison-based** sorts.

(They compare things, using some ordering relation  $\leq$ .)

Works on *any* data, so long as we have  $\leq$ .

What if we know more about the keys?  
... could we get down to  $O(n)$  time?

# Key-Indexed Counting Sort

count up the number of times each key appears;  
this indexes where each item belongs in the sorted array

FOR EXAMPLE:

assuming my key domain is numbers  $[0 \dots 10]$ ,  
if we have three '0's, and two '1's,  
'2's must go at index 5 and onwards

look, ma! no comparisons!

look, ma! **an  $O(n)$  sort!**

terms and conditions apply, see in store for details

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

we must know our sequence is of size  $N$ ,  
and the domain of keys in that sequence.

pumped-up KICS pretty efficient  
... if  $M$  is small compared to  $N$ .  
actually,  $O(n + M)$  ... so if we have 1, 2, 999999 ...

Not in-place — uses a temporary array.  
Can be stable! Not really adaptive.

We already have a data structure  
which has element ordering as an invariant:  
the **heap** or **priority queue**.

We *could* just dump all  $n$  elements  
into a priority queue, and dequeue them —  
 $n$  operations of  $O(\log n)$  complexity.  
no gain.

What if we used the heap-fix-down mechanism  
on the whole array, popping off the maximum item,  
and shrinking the heap each time? That's  $O(n)$ !  
The catch: the inner loop is expensive.

## Sorting

Divide-and-Conquer

Merge

Quick

Non-Comparison

Key-Indexed

Heap

```
void sort_heap (Item a[], size_t lo, size_t hi)
{
    size_t N = hi - lo + 1;
    Item *pq = &a[lo - 1];
    for (size_t k = N/2; k >= 1; k--)
        heap_fixdown (pq, k, N);
    while (N > 1) {
        swap_idx (pq, 1, N);
        heap_fixdown (pq, 1, --N);
    }
}
```